

AD 6 5 4 : Marketing Analytics

Boston University

Assignment 1 : Data Exploration & Visualization

For **Parts I & II** of this assignment, you will use the file **lobster 2 4 .csv**, which can be found in Blackboard on our course discussion forum under Class Discussion >> From Your Professor. The dataset includes information about Lobster Land operations throughout the **2 0 2 4** seas which began on Memorial Day of this year, and concluded on Labor Day of this year. Part III of the assignment will use the **lobsterland_ 2 0 2 4 .csv** file, which can also be found in Blackboard course discussion forum under Class Discussion >> From Your Professor.

Once you have completed this assignment, you will upload three files into Blackboard:

- For **Parts I, II, and III** of this assignment:
 - The .ipynb file that you create in Jupyter Notebook (or Colab)
 - Include a one-paragraph description of your dashboard for Part III.
 - A PDF file that was generated from your .ipynb file
- For **Part III** of this assignment:
 - Either your .twb file (if you used Tableau Desktop) or a link to your dashboard on the Tableau server (if you used Tableau Public)

The .ipynb file that you create in Jupyter Notebook (or Colab), and a PDF file that was generated from your .ipynb file. If you are using Jupyter Notebook, you may wish to download your work as an .html and then save that file as a PDF. **Please include your last name in your filenames.** The exact name it is up to you, but the last name makes it easier to keep track of the file (e.g. BakerAssignment 1.ipynb, bakerAssgn 1.ipynb, bakerAssignment 1.html, etc. -- any of these will do). Blackboard will not accept an .html file, which is why you should upload a PDF and an ipynb.

For any question that asks you to perform some particular task, you just need to show your input output. Tasks will always be written in regular, non-italicized font.

For any question that asks you to include interpretation, write your answer in a Markdown cell in Notebook (or a 'Text' cell if you use Colab). *Any homework question that needs interpretation will be written in italicized font.* Do not simply write your answer in a code cell as a comment, but use a Markdown or Text cell instead.

Remember to be resourceful! There are many helpful resources available to you, including the vi library, the lecture notes on Blackboard, facilitator sessions, the office hours sessions, and the w

In the prompt, variables might not be referred to in the exact way that their names appear in the code. This is okay -- that's very realistic. You should familiarize yourself with the dataset and its variables through the dataset description table.

Dataset Description:

Variable	Description
Date	The date, in year-month-date format. Note that each row in the dataset corresponds to one day of park operations.
Total_Visitors	The total number of visitors that came to Lobster Land on that particular day.
Arcade_Visitors	The total number of visitors who came to the Gold Zone (the Lobster Land arcade) on that particular day.
Total_Purchases	The total amount of spending by Lobster Land visitors on the particular day.
Arcade_Revenue	The total amount of revenue generated by Gold Zone spending on the particular day.
Total_Labor_Hours	The total amount of labor hours on the particular day. Each park employee is paid \$15/hour. If one employee works a 6-hour shift, this contributes 6 total labor hours to the statistic.
Weather	Type A categorical descriptor for the type of weather for the day.
Special_Events	This indicates what type of special event (if any) was held that day. Note that days without special events show a value of 'NaN' in this column.
Customer_Complaints	This indicates the total number of customer complaints, as registered by the Lobster Land Customer Service Desk, on a particular day of park operations.

|Day of Week| This indicates the day of the week. The days are mostly distributed in an even way since the season both starts and ends on Mondays, there is one more Monday (compared to the other days of the week) |Passholder Percentage| This indicates the percentage of visitors per day who are Lobster Land passholders. |International Visitors| This indicates the number of visitors per day from international visitors. |High Temperature| This is the high temperature per day| |General Weather| This is a categorical descriptor for the type of weather for the day.| |Precipitation| This column takes a value of 1 or a 0 value to indicate whether any precipitation occurred during that 2 4 -hour period.

Your Tasks:

Bring the **Lobster 2 4 .csv** dataset into your local environment (in Jupyter Notebook, or in Colab)

Part I - Exploratory Data Analysis: Exploration Manipulation

A. Call the `head()` function on this dataframe and look at your results.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set display options
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)

# Load the dataset
df = pd.read_csv('lobster24.csv')
print("Dataset loaded successfully!")
print(f"Shape: {df.shape}")
df.head()
```

Dataset loaded successfully!

Shape: (99, 15)

```
Out[2]:
```

	Date	Total_Visitors	Arcade_Visitors	Total_Purchases	Arcade_Reven
0	5 / 2 7 / 2 0 2 4	1 0 6 0	2 6 7	2 2 5 7 5 . 7 1	3 5 2 4 . 3
1	5 / 2 8 / 2 0 2 4	1 4 9 4	5 5 2	1 7 1 4 2 . 5 2	1 6 5 2 . 9
2	5 / 2 9 / 2 0 2 4	1 3 3 0	4 4 7	1 0 2 2 9 . 9 9	3 2 5 1 . 8
3	5 / 3 0 / 2 0 2 4	1 2 9 5	4 4 2	6 4 4 2 . 6 5	8 6 7 . 1
4	5 / 3 1 / 2 0 2 4	1 8 3 8	2 5 6	2 8 4 0 9 . 0 6	7 4 7 . 6

B. How many rows of the dataset are visible in Jupyter now?

There are 5 visible rows

C. Take a look at the dataset's shape attribute.

- a. How many rows, and how many columns, are in this entire dataframe?

```
In [3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99 entries, 0 to 98
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date             99 non-null      object  
 1   Total_Visitors   99 non-null      int64   
 2   Arcade_Visitors  99 non-null      int64   
 3   Total_Purchases  99 non-null      float64 
 4   Arcade_Revenue   99 non-null      float64 
 5   Total_Labor_Hours 99 non-null      float64 
 6   Weather_Type     99 non-null      object  
 7   Special_Events   21 non-null      object  
 8   Customer_Complaints 99 non-null      int64   
 9   Day_of_Week       99 non-null      object  
 10  Passholder_Percentage 99 non-null      float64 
 11  International_Visitors 99 non-null      int64   
 12  High_Temperature 96 non-null      float64 
 13  General_Weather   99 non-null      object  
 14  Precipitation    99 non-null      int64   
dtypes: float64(5), int64(5), object(5)
memory usage: 11.7+ KB

```

There are 9 9 rows and 1 5 columns in the dataframe

D. Read the dataset description, and take a look at the variables in the dataset.

- a. *Which of your variables should be seen as categorical, and which ones should be seen as numeric?*

CATEGORICAL VARIABLES:

- Date
- Weather_Type
- Special_Events
- Day_of_Week
- General_Weather

NUMERIC VARIABLES:

- Total_Visitors
- Arcade_Visitors
- Total_Purchases
- Arcade_Revenue
- Total_Labor_Hours
- Customer_Complaints
- Passholder_Percentage
- International_Visitors
- High_Temperature
- Precipitation

- a. Are there any `NaN` values in this dataset? Show the code that you used to answer this question. Which columns have `NaN` values?

```
In [4]: missing = df.isna().sum().sort_values(ascending=False)
missing[missing > 0]
```

```
Out[4]: Special_Events      78
High_Temperature       3
dtype: int64
```

There are missing values in `Special_Events` (78) and `High_Temperature` (3)

- b. Take a close look at the dataset description. Is there a variable in this dataset that contains values, but for which the dataset could be more clearly descriptive / informative about the actual values? If so, change the '`NaN`' values in that column to something that's more descriptive.

```
In [5]: df.describe()
```

	Total_Visitors	Arcade_Visitors	Total_Purchases
count	9 9.0 0 0 0 0 0 0	9 9.0 0 0 0 0 0 0	9 9.0 0 0 0 0 0 0
mean	1 3 9 9.7 3 7 3 7 4	4 0 8 0 5 0 5 0 5	1 5 7 1 2 0 1 1 8 1 8
std	6 5 1.1 5 7 8 0 1	1 8 2.4 6 0 1 3 7	5 7 3 7.1 6 3 1 0 2
min	2 2 1.0 0 0 0 0 0	5 1.0 0 0 0 0 0 0	5 2 7 1.1 3 0 0 0 0
2 5 %	9 2 4.5 0 0 0 0 0	2 6 8.0 0 0 0 0 0	1 1 0 6 9.2 7 5 0 0 0
5 0 %	1 3 8 0.0 0 0 0 0	4 2 6.0 0 0 0 0 0	1 6 0 2 8.2 4 0 0 0 0
7 5 %	1 9 2 4.5 0 0 0 0	5 5 9.5 0 0 0 0 0	1 9 0 0 3.9 7 0 0 0 0
max	2 4 7 8.0 0 0 0 0 0	6 9 8.0 0 0 0 0 0	2 8 5 9 9.9 6 0 0 0 0

- c. Is there a column in the dataset with some `NaN` values, for which the true values could be found? (keep in mind that Lobster Land is based just outside of Portland, Maine). If so, find the actual values that are missing, and replace them in the dataset. Which values did you use for `NaN` replacement?

```
In [6]: df['Special_Events'] = df['Special_Events'].fillna('No Special Event')

print("\nAfter change - Special_Events value counts:")
print(df['Special_Events'].value_counts(dropna=False))
```

After change - `Special_Events` value counts:

Special_Events	count
No Special Event	78
Arcade Tournament	6
Fireworks	6
Food Festival	5
Music Festival	4
Name: count, dtype: int64	

The Special_Events column had NaN values for days without special events and instead of showing 'NaN', I replaced them with 'No Special Event'

```
In [7]: print("\nRows with missing High_Temperature:")
missing_temp = df[df['High_Temperature'].isnull()]
print(missing_temp[['Date', 'High_Temperature', 'General_Weather']])
```

```
Rows with missing High_Temperature:
      Date  High_Temperature General_Weather
22   6/18/2024           NaN    Rainy/Stormy
44   7/10/2024           NaN    Rainy/Stormy
67   8/2/2024            NaN       Mild
```

Actual temperatures from Portland, Maine weather records

Source: Extreme Weather Watch (NOAA data)

<https://www.extremeweatherwatch.com/cities/portland-me/year-2024>

```
In [8]: df.loc[df['Date'] == '6/18/2024', 'High_Temperature'] = 85.0
df.loc[df['Date'] == '7/10/2024', 'High_Temperature'] = 87.0
df.loc[df['Date'] == '8/2/2024', 'High_Temperature'] = 88.0

print("\nMissing temperatures filled with actual values from Extreme Weather Watch")
print(" - 6/18/2024: 85°F")
print(" - 7/10/2024: 87°F")
print(" - 8/2/2024: 88°F")

# Verify the fill
print("\nVerification:")
print(df[df['Date'].isin(['6/18/2024', '7/10/2024', '8/2/2024'])][['Date']]
```

Missing temperatures filled with actual values from Extreme Weather Watch (NO data):

- 6/18/2024: 85°F
- 7/10/2024: 87°F
- 8/2/2024: 88°F

Verification:

	Date	High_Temperature	Weather_Type
22	6/18/2024	85.0	Rainy
44	7/10/2024	87.0	Stormy
67	8/2/2024	88.0	Sunny

The High_Temperature column had 3 missing values. Since Lobster Land is located near Farnsworth Beach, I looked up the actual historical temperature data from Extreme Weather Watch (<https://www.extremeweatherwatch.com/cities/portland-me/year-2024>)

www.extremeweatherwatch.com/cities/portland-me/year-2024), which uses NOAA data, for Portland, Maine on these specific dates:

- June 18, 2024: 85°F
- July 10, 2024: 87°F
- August 2, 2024: 88°F

These values have been filled into the dataset.

F. Lobster Land management just realized that they want to change something in their dataset! They think it would make more sense to see the day of the week column appear next to the 'Date' column. Right now, the 'Date' column is to the far left of the dataframe, but the day of the week is somewhere else.

- a. Use some Python code to move the day of the week column to put it immediately to the right of the 'Date' column.

```
In [9]: cols = df.columns.tolist()
cols.remove('Day_of_Week')
date_index = cols.index('Date')
cols.insert(date_index + 1, 'Day_of_Week')
df = df[cols]

print("\nNew column order (first 5):")
print(list(df.columns[:5]))
```

New column order (first 5):
['Date', 'Day_of_Week', 'Total_Visitors', 'Arcade_Visitors', 'Total_Purchases']

- b. Next, rename the 'Total_Purchases' variable to instead be named 'Total_Revenue.'

```
In [10]: df = df.rename(columns={'Total_Purchases': 'Total_Revenue'})

print("\nRenamed Total_Purchases to Total_Revenue")
print("Revenue columns:", [col for col in df.columns if 'Revenue' in col])
```

Renamed Total_Purchases to Total_Revenue
Revenue columns: ['Total_Revenue', 'Arcade_Revenue']

G. Lobster Land wants to better understand the way in which management allocates hours for their workers. Specifically, management wants to know about efficiency. Create a new variable called 'staff_efficiency'. This variable should be found by taking total revenue and dividing by total staff |

```
In [11]: df['staff_efficiency'] = df['Total_Revenue'] / df['Total_Labor_Hours']

print("Sample of staff_efficiency values:")
print(df[['Date', 'Total_Revenue', 'Total_Labor_Hours', 'staff_efficiency']])
```

Sample of staff_efficiency values:

	Date	Total_Revenue	Total_Labor_Hours	staff_efficiency
0	5/27/2024	22575.71	185.84	121.479283
1	5/28/2024	17142.52	250.90	68.324113
2	5/29/2024	10229.99	115.59	88.502379
3	5/30/2024	6442.65	347.30	18.550677
4	5/31/2024	28409.06	234.62	121.085415

- a. Now that you have created this variable, sort the dataset by staff_efficiency, from highest to lowest.

```
In [12]: # Sort by efficiency (highest to lowest)
df_sorted = df.sort_values('staff_efficiency', ascending=False)
print("\nDataset sorted by staff_efficiency")
```

Dataset sorted by staff_efficiency

- b. Isolate the 10 most efficient days from the 2024 season, and take a look at them. *What can you say about these 10 days? Do they look different from the overall or do they instead look like a random sample of the rest of the days?* Answer this in 3 - 5 sentences, and use any evidence from the dataset as you see fit.

```
In [13]: top_10_efficient = df_sorted.head(10)
print("\nTop 10 Most Efficient Days:")
print(top_10_efficient[['Date', 'Day_of_Week', 'Weather_Type', 'Special_Even
                                'Total_Revenue', 'Total_Labor_Hours', 'staff_efficiency']])
```

Top 10 Most Efficient Days:

	Date	Day_of_Week	Weather_Type	Special_Events	Total_Revenue	staff_efficiency
36	7/2/2024	Tuesday	Rainy	Fireworks	19604.96	140.045432
97	9/1/2024	Sunday	Sunny	Food Festival	18142.97	136.249399
0	5/27/2024	Monday	Sunny	Arcade Tournament	22575.71	121.479283
4	5/31/2024	Friday	Sunny	Arcade Tournament	28409.06	121.085415
69	8/4/2024	Sunday	Stormy	No Special Event	17466.68	110.102622
46	7/12/2024	Friday	Sunny	No Special Event	14499.29	106.082016
56	7/22/2024	Monday	Stormy	Fireworks	12301.13	100.540499
52	7/18/2024	Thursday	Stormy	No Special Event	14544.99	100.041200
18	6/14/2024	Friday	Rainy	Music Festival	26427.39	98.100857
30	6/26/2024	Wednesday	Cloudy	No Special Event	19034.52	97.467971

	Total_Labor_Hours	staff_efficiency
36	139.99	140.045432
97	133.16	136.249399
0	185.84	121.479283
4	234.62	121.085415
69	158.64	110.102622
46	136.68	106.082016
56	122.35	100.540499
52	145.39	100.041200
18	269.39	98.100857
30	195.29	97.467971

The top 10 most efficient days reveal clear differences from the broader dataset. Many of high efficiency days coincide with special events especially Music Festivals, Fireworks shows, and Festivals, which generate strong revenue without a proportional rise in labor hours. Weather can vary across these days, indicating that well designed events can offset less than ideal weather. E

is primarily driven by high revenue levels (2 8 K) paired with tightly managed labor hours (below 3 0 0), suggesting either intentional staffing optimization or naturally high demand. Overall, this highlights strategic event planning as a powerful lever for boosting operational efficiency.

H. Create a new variable called 'international_percentage'. This variable should indicate the percentage of all visitors who were international, on any given day (so for instance, this would show ' 5 0 ' if there were 4 0 0 0 total visitors on a given day 2 0 0 0 of them were international).

```
In [14]: # Create international_percentage variable
df['international_percentage'] = (df['International_Visitors'] / df['Total_Visitors']) * 100

print("Sample of international_percentage:")
print(df[['Date', 'Total_Visitors', 'International_Visitors', 'international_percentage']])
```

Sample of international_percentage:

	Date	Total_Visitors	International_Visitors	international_percentag
0	5/27/2024	1060	357	33.67924
1	5/28/2024	1494	424	28.38018
2	5/29/2024	1330	441	33.15789
3	5/30/2024	1295	338	26.10038
4	5/31/2024	1838	615	33.46028
5	6/1/2024	2369	859	36.26002
6	6/2/2024	666	259	38.88888
7	6/3/2024	1438	474	32.96244
8	6/4/2024	530	140	26.41509
9	6/5/2024	1682	459	27.28894

- a. Next, find the correlation between international percentage and passholder percentage. *What is this correlation? Is it positive or negative? Strong, moderate, or weak? Next, take a moment to reason as to 'why' this relationship might look the way it does – what could explain it?*

```
In [15]: # First convert Passholder_Percentage to actual percentage (multiply by 100)
correlation = df['international_percentage'].corr(df['Passholder_Percentage'])

print(f"\nCorrelation: {correlation:.4f}")

# Interpret correlation
if correlation > 0:
    strength_direction = "positive"
else:
    strength_direction = "negative"

abs_corr = abs(correlation)
if abs_corr > 0.7:
    strength = "strong"
elif abs_corr > 0.4:
    strength = "moderate"
else:
    strength = "weak"

print(f"This correlation is {strength_direction} and {strength}.")
```

Correlation: -0.6722
This correlation is negative and moderate.

The correlation of - 0 . 6 7 indicates a strong negative relationship. As international visitors ir passholder percentage tends to decrease. This makes sense because passholders are mostly local repeat visitors, while international visitors are typically tourists who come during peak periods. As a result, days with more tourists tend to have fewer locals, and vice versa. These patterns reflect two distinct segments: locals optimizing for comfort and flexibility, and tourists bound by fixed travel schedules.

I. Use the `info()` function to determine the data type for the 'Date' variable. *What type of variable does pandas see it as, currently?* Convert this variable into a datetime object, and show that the conversion was successful.

```
In [16]: # Check current data type
print("Current data type of 'Date' column:")
print(f"Date is type: {df['Date'].dtype}")

# Convert to datetime
df['Date'] = pd.to_datetime(df['Date'])

print("\nAfter conversion:")
print(f"Date is now type: {df['Date'].dtype}")
print("\nSample of Date column:")
print(df['Date'].head())
```

Current data type of 'Date' column:
Date is type: object

After conversion:
Date is now type: datetime64[ns]

Sample of Date column:
0 2024-05-27
1 2024-05-28
2 2024-05-29
3 2024-05-30
4 2024-05-31
Name: Date, dtype: datetime64[ns]

- a. What impact could this have on the variable? (In other words, what is the point of having Python 'see' that your date variable is truly a date?)

Converting a Date column into a true datetime object lets Python treat it as an actual date instead of text. This unlocks easy extraction of things like month or weekday, enables time-series analysis (date differences, rolling averages), and ensures accurate sorting and filtering by date. It also improves readability since visualization tools can format time-based axes correctly.

J. Using the `groupby()` function, along with the `describe()` function, explore the relationship between day of the week and total revenue. *What stands out here? In which sentences, describe anything noteworthy, and also include some speculation about why the results look the way that they do.*

```
In [17]: revenue_by_day = df.groupby('Day_of_Week')['Total_Revenue'].describe()
print("Revenue statistics by day of week:")
```

```
print(revenue_by_day)
print("\n")
```

Revenue statistics by day of week:

Day_of_Week	count	mean	std	min	25%	\
Friday	14.0	17133.101429	6309.542051	8052.76	12365.9700	
Monday	15.0	13525.969333	5161.017456	7624.32	9002.9300	
Saturday	14.0	21448.985000	4161.667899	15353.92	18287.3075	
Sunday	14.0	18807.897143	4946.299831	11984.34	15227.6475	
Thursday	14.0	12125.405000	4653.834452	6056.24	7931.7700	
Tuesday	14.0	14155.262143	5128.831358	5271.13	10193.9950	
Wednesday	14.0	12943.608571	3422.424946	9183.07	9936.1925	
Day_of_Week		50%	75%	max		
Friday	15628.680	21562.4450	28409.06			
Monday	12301.130	18278.0500	22575.71			
Saturday	21655.475	23664.5925	28599.96			
Sunday	17804.825	21440.6825	28132.69			
Thursday	12138.675	16272.5750	19296.08			
Tuesday	16138.935	18130.2900	19732.53			
Wednesday	12520.310	15416.1950	19034.52			

Saturday delivers the highest revenue, reflecting classic weekend demand when families and tourists have more free time. Sunday is similarly strong. In contrast, Wednesday and Thursday are the slowest days, making them ideal for lighter staffing, maintenance, or targeted mid-week promotions. Friday shows a clear ramp-up as visitors begin their weekend plans.

Weekend revenue also varies more from week to week, likely because weather and special events have a bigger impact when baseline attendance is already high.

K. *Which two variables in this dataset are not identical, but mostly redundant?* Identify this pair of variables, and remove one of them (but not both!) You will be able to complete all of the subsequent steps, regardless of which one you remove here.

In [18]:

```
print("Examining potentially redundant variables...")
print("\nWeather-related columns:")
print(df[['Weather_Type', 'General_Weather', 'Precipitation']].head(15))

print("\nUnique values in Weather_Type:")
print(df['Weather_Type'].value_counts())

print("\nUnique values in General_Weather:")
print(df['General_Weather'].value_counts())

print("\nCross-tabulation of Weather_Type vs General_Weather:")
print(pd.crosstab(df['Weather_Type'], df['General_Weather']))
```

Examining potentially redundant variables...

Weather-related columns:

	Weather_Type	General_Weather	Precipitation
0	Sunny	Cool	0
1	Cloudy	Cool	0
2	Cloudy	Hot	0
3	Stormy	Rainy/Stormy	1
4	Sunny	Mild	0
5	Stormy	Rainy/Stormy	1
6	Sunny	Mild	0
7	Sunny	Mild	0
8	Rainy	Rainy/Stormy	1
9	Sunny	Cool	0
10	Stormy	Rainy/Stormy	1
11	Cloudy	Mild	0
12	Sunny	Cool	0
13	Rainy	Rainy/Stormy	1
14	Stormy	Rainy/Stormy	1

Unique values in Weather_Type:

```
Weather_Type
Sunny      30
Stormy     29
Cloudy     23
Rainy      17
Name: count, dtype: int64
```

Unique values in General_Weather:

```
General_Weather
Rainy/Stormy    46
Mild            23
Cool             21
Hot              9
Name: count, dtype: int64
```

Cross-tabulation of Weather_Type vs General_Weather:

General_Weather	Cool	Hot	Mild	Rainy/Stormy
Weather_Type				
Cloudy	9	4	10	0
Rainy	0	0	0	17
Stormy	0	0	0	29
Sunny	12	5	13	0

Weather_Type and General_Weather are essentially redundant because they describe the same underlying conditions. Weather_Type focuses on the core attendance drivers sun, clouds, rain, a storms while General_Weather blends temperature (Cool/Mild/Hot) with precipitation. Since temp and precipitation are already captured more precisely in High_Temperature and Precipitation, General_Weather doesn't add meaningful new information.

Weather_Type is more actionable and directly tied to operational decisions, making it the better \ to keep.

Decision: Drop General_Weather and retain Weather_Type.

```
In [19]: # Remove the redundant variable
df = df.drop('General_Weather', axis=1)
```

```
print("\nGeneral_Weather column removed.")
print(f"New dataset shape: {df.shape}")
print(f"Now have {df.shape[1]} columns instead of 15")
```

```
General_Weather column removed.
New dataset shape: (99, 16)
Now have 16 columns instead of 15
```

Part II - Data Visualization

L. Using any plotting tool in Python, generate a bar plot that shows the weather type on one axis, and shows average (mean) total revenue on the other axis. Be sure to give your plot a clear, descriptive title.

```
In [20]: # Group by Weather_Type and calculate mean revenue
weather_rev = df.groupby('Weather_Type')[['Total_Revenue']].mean().reset_index()


# Create bar plot
plt.figure(figsize=(8,5))
plt.bar(weather_rev['Weather_Type'], weather_rev['Total_Revenue'], color='blue')

# Title and labels
plt.title('Average Total Revenue by Weather Type')
plt.xlabel('Weather Type')
plt.ylabel('Average Total Revenue ($)')
plt.xticks(rotation=45)

# Add value labels on top of bars
for i, (weather, value) in enumerate(weather_rev.iterrows()):
    plt.text(i, value['Total_Revenue'] + 200, f'${value["Total_Revenue"]}', ha='center', fontweight='bold')

# Sort bars by revenue (highest to lowest)
weather_rev = weather_rev.sort_values('Total_Revenue', ascending=False)

plt.tight_layout()
plt.show()
```



- a. *What do you notice about this plot? Does it surprise you, or does it align with your expectations? In a couple of sentences, how would you describe what you see here?*


Comparing 10 bins to 30 bins makes the impact of bin size much clearer. With 10 bins, the histogram looks smoother and hides many of the smaller fluctuations in arcade revenue. Increasing the number of bins to 30 reveals more detail, showing small peaks that were not visible before. This illustrates how larger bin counts can expose day-to-day variability, while smaller bin counts smooth the data into broader patterns.

M. Next, generate a histogram that depicts total spending at the Gold Zone (the Lobster Land arcade) across the days of the 2024 season. Be sure to give your plot a clear, descriptive title.

```
In [21]: plt.figure(figsize=(10, 6))
plt.hist(df['Arcade_Revenue'], bins=10, color="#06A77D", edgecolor='black')
plt.xlabel('Arcade Revenue ($)', fontsize=12, fontweight='bold')
plt.ylabel('Frequency (Number of Days)', fontsize=12, fontweight='bold')
plt.title('Distribution of Daily Arcade Revenue (Gold Zone)\nLobster Land', fontsize=14, fontweight='bold', pad=20)
plt.grid(axis='y', alpha=0.3)

plt.tight_layout()
plt.show()
```


Distribution of Daily Arcade Revenue (Gold Zone) Lobster Land 2024 Season - 10 Bins



```
In [22]: plt.figure(figsize=(10, 6))
n, bins, patches = plt.hist(df['Arcade_Revenue'], bins=30,
                             color='#06A77D', edgecolor='black', alpha=.75)

# Add count labels on top of each bin
for i in range(len(patches)):
    height = patches[i].get_height()
    if height > 0: # Only label non-zero bars
        plt.text(patches[i].get_x() + patches[i].get_width()/2, height +
                  f'{int(height)}', # Shows count
                  ha='center', va='bottom', fontweight='bold', fontsize=9)

plt.xlabel('Arcade Revenue ($)')
plt.ylabel('Frequency (Number of Days)')
plt.title('Distribution of Daily Arcade Revenue')
plt.show()
```



- a. In a couple of sentences, describe your plot. What does it show you about the distribution of arcade spending in the 2024 season?

Comparing 10 bins to 30 bins makes the impact of bin size much clearer. With 10 bins, the histogram looks smoother and hides many of the smaller fluctuations in arcade revenue. Increasing to 30 bins, the distribution becomes much more detailed, revealing small peaks that were not visible before. This shows how larger bin counts can expose day-to-day variability, while smaller bin counts smooth the data into broader patterns.


- b. Now, increase the number of bins. What can you say about this second histogram, compared to the first one? What impact can changing the number of bins (either higher or lower) have on a histogram?

```
In [23]: plt.figure(figsize=(10, 6))
n, bins, patches = plt.hist(df['Arcade_Revenue'], bins=13,
                             color='#06A77D', edgecolor='black', alpha=0.8)

# Add count labels on top of each bin
for i in range(len(patches)):
    height = patches[i].get_height()
    if height > 0: # Only label non-zero bars
        plt.text(patches[i].get_x() + patches[i].get_width()/2, height +
                  f'{int(height)}', # Shows count
                  ha='center', va='bottom', fontweight='bold', fontsize=9)

plt.xlabel('Arcade Revenue ($)')
plt.ylabel('Frequency (Number of Days)')
```

```
plt.title('Distribution of Daily Arcade Revenue')
plt.show()
```




The `1 3`-bin histogram reveals more small clusters than the `1 0`-bin version, showing that revenue varies across many levels rather than just “low vs. high.” This added detail makes the plot less smooth, so fewer bins are better for big-picture storytelling while more bins are useful for detailed analytical work.

N. One member of the Lobster Land Board of Directors thinks that people tend to complain more on hot days. Let’s see whether the data supports this theory! Generate a scatterplot with the day’s high temperature on the x-axis, and with the number of customer complaints on the y-axis. Be sure to give your plot a clear, descriptive title.

```
In [24]: plt.figure(figsize=(10, 6))
plt.scatter(df['High_Temperature'], df['Customer_Complaints'],
            alpha=0.6, s=100, color='#F77F00', edgecolor='black', linewidth=1)
plt.xlabel('High Temperature (°F)', fontsize=12, fontweight='bold')
plt.ylabel('Number of Customer Complaints', fontsize=12, fontweight='bold')
plt.title('Customer Complaints vs Daily High Temperature\nLobster Land 2024 Season',
          fontsize=14, fontweight='bold', pad=20)
plt.grid(True, alpha=0.3)

# Add a trend line
z = np.polyfit(df['High_Temperature'], df['Customer_Complaints'], 1)
p = np.poly1d(z)
plt.plot(df['High_Temperature'].sort_values(),
         p(df['High_Temperature'].sort_values()),
         "r--", alpha=0.8, linewidth=2, label='Trend Line')
plt.legend()

plt.tight_layout()
plt.show()
```



- a. What does the scatterplot show? If there is a relationship between these variables, what can you explain it? If there is not a relationship between these variables, why do you think this is the case?

The scatterplot indicates almost no relationship between temperature and customer complaints (correlation $\approx \{0.1\}$); the points are widely dispersed, showing that hotter days cannot reliably produce more complaints than cooler ones. This directly challenges the Board member's assumption that heat drives dissatisfaction. Several factors could explain the lack of a temperature effect:

1. Comfort Management: Cooling stations, shade, and water features may effectively offset heat discomfort.

- 2 . Operational Drivers: Complaints likely stem from controllable issues such as wait times, staff interactions, or food quality rather than weather.
- 3 . Behavioral Adjustment: On very hot days, guests may modify their activities or shorten visits of complaining.
- 4 . Expectation Setting: Summer visitors anticipate heat and arrive prepared, making them less sensitive to temperature than to operational problems.

Operationally, this points toward prioritizing improvements in service quality and guest experience over expanding weather-related amenities.


O. Generate a boxplot that shows day of week on one axis, and passholder perc on the other axis. Be sure to give your plot a clear, descriptive title.

```
In [25]: # Order days of week properly
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
df['Day_of_Week'] = pd.Categorical(df['Day_of_Week'], categories=day_order)

plt.figure(figsize=(12, 6))
sns.boxplot(data=df, x='Day_of_Week', y='Passholder_Percentage', linewidth=2)

plt.xlabel('Day of Week', fontsize=12, fontweight='bold')
plt.ylabel('Passholder Percentage', fontsize=12, fontweight='bold')
plt.title('Passholder Percentage by Day of Week\nLobster Land 2024 Season', fontsize=14, fontweight='bold', pad=20)
plt.xticks(rotation=45)
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()

print("\nPassholder Percentage by Day of Week:")
df.groupby('Day_of_Week', observed=False)[['Passholder_Percentage']].describe()
```



Passholder Percentage by Day of Week:

Out[25]:

	count	mean	std	min	2 5 %
Day_of_Week					
Monday	1 5 .0	0 .2 5 1 3 3 3	0 .0 3 8 8 8 9	0 .1 8	0 .2 2 0 0 (
Tuesday	1 4 .0	0 .2 2 9 2 8 6	0 .0 3 6 6 8 2	0 .1 6	0 .2 1 0 0 (
Wednesday	1 4 .0	0 .2 3 0 7 1 4	0 .0 4 1 0 3 7	0 .1 6	0 .2 0 2 5 (
Thursday	1 4 .0	0 .2 4 7 1 4 3	0 .0 3 9 6 9 7	0 .1 8	0 .2 3 2 5 (
Friday	1 4 .0	0 .2 3 5 0 0 0	0 .0 3 8 9 7 7	0 .1 7	0 .2 1 0 0 (
Saturday	1 4 .0	0 .1 8 4 2 8 6	0 .0 2 7 3 7 6	0 .1 5	0 .1 6 0 0 (
Sunday	1 4 .0	0 .2 0 8 5 7 1	0 .0 2 5 6 7 8	0 .1 7	0 .1 8 2 5 (

- a. In a few sentences, what does this plot show you? Why do you think it looks this way – what explains this?

The box plot shows that passholder percentages are higher on weekdays and lower on weekends. The weekday values, especially Monday, Tuesday, and Wednesday, also vary much more from day to day. Their taller boxes and larger standard deviations indicate that these days are less predictable, while Saturday and Sunday stay in a tight, stable range even though their percentages are lower. This suggests that weekday attendance is influenced by more fluctuating factors (work schedules, school calendar, local behavior), whereas weekend patterns are more consistent.

P. Generate a barplot that shows Special Events on one axis, and international visitors on the other axis. Do not include confidence intervals ('error bars') on your plot. Order your bars either from shortest to tallest, or from tallest to shortest. Be sure to give your plot a clear, descriptive title.

In [26]:


```
# Calculate mean international visitors by event type
event_intl = df.groupby('Special_Events')['International_Visitors'].mean()

plt.figure(figsize=(12, 6))
plt.bar(range(len(event_intl)), event_intl.values, color='#023047', edgecolor='black')
plt.xlabel('Special Event Type', fontsize=12, fontweight='bold')
plt.ylabel('Average International Visitors', fontsize=12, fontweight='bold')
plt.title('Average International Visitors by Special Event Type\nLobsterfest', fontsize=14, fontweight='bold', pad=20)
plt.xticks(range(len(event_intl)), event_intl.index, rotation=45, ha='right')
plt.grid(axis='y', alpha=0.3)

# Add value labels
for i, value in enumerate(event_intl.values):
    plt.text(i, value + 20, f'{value:.0f}', ha='center', va='bottom', fontweight='bold', fontsize=10)

plt.tight_layout()
plt.show()

print("\nAverage International Visitors by Event Type:")
print(event_intl)
```



Average International Visitors by Event Type:

`Special_Events`

Fireworks	601.333333
Music Festival	530.500000
No Special Event	494.858974
Food Festival	449.400000
Arcade Tournament	447.833333

`Name: International_Visitors, dtype: float64`

- a. What does your plot show? Why can we **not** conclude causality based on the relationship here?

Fireworks events draw the highest average international attendance, suggesting they're a major attraction. However, the "No Special Event" category still shows strong visitor numbers, indicating international guests attend Lobster Land consistently, not just for special events.

Q. Create a faceted bar plot showing arcade revenue, with facets for different weather types and bars representing each day of the week. Do not show confidence intervals ('error bars') on your plot. Be sure to give your plot a clear, descriptive title.

```
In [27]: # Create faceted plot
weather_types = df['Weather_Type'].unique()
fig, axes = plt.subplots(2, 2, figsize=(16, 12))
axes = axes.flatten()

for idx, weather in enumerate(weather_types):
    weather_data = df[df['Weather_Type'] == weather]
    day_arcade = (
        weather_data
        .groupby('Day_of_Week', observed=True)[ 'Arcade_Revenue']
        .mean()
        .reindex(day_order)
    )

    axes[idx].bar(
        range(len(day_arcade)),
        day_arcade.values,
```

```

        color=[  

            '#1B4965', '#5FA8D3', '#62B6CB', '#CAE9FF',  

            '#FFB703', '#FB8500', '#DC2F02'  

        ][:len(day_arcade)]  

    )  

    axes[idx].set_title(f'{weather} Days', fontsize=14, fontweight='bold')  

    axes[idx].set_xlabel('Day of Week', fontsize=11, fontweight='bold')  

    axes[idx].set_ylabel('Avg Arcade Revenue ($)', fontsize=11, fontweight='bold')  

    axes[idx].set_xticks(range(len(day_arcade)))  

    axes[idx].set_xticklabels(day_arcade.index, rotation=45, ha='right')  

    axes[idx].grid(axis='y', alpha=0.3)  

plt.suptitle('Average Arcade Revenue by Weather Type and Day of Week\nLobster Land 2024 Season',  


             fontsize=16, fontweight='bold', y=1.00)  

plt.tight_layout()  

plt.show()

```

Average Arcade Revenue by Weather Type and Day of Week
Lobster Land 2024 Season



In [28]: `df['Weather_Type'].value_counts()`

Out[28]: Weather_Type

Sunny	30
Stormy	29
Cloudy	23
Rainy	17

Name: count, dtype: int64

- a. What looks interesting here? Are there some day of week/weather combinations that seem particularly noteworthy? Also, what's an important limitation to note when interpreting this graph?

The plot shows that sunny days consistently generate the highest arcade revenue, especially on Mondays and Thursdays, which stand out across multiple weather types. Cloudy days perform reasonably well, while stormy and rainy weekends tend to show the lowest revenue. One interesting pattern is that weekday revenue remains relatively strong even in poor weather, suggesting that loyal passholders may still visit the arcade regardless of conditions.

Important limitation A key limitation is that some day-of-week and weather-type combinations are rare or have very few observations. For example, we may have many Sunny Fridays but only one Stormy Wednesday. This means the averages in the plot are not based on equal sample sizes, so some represent much more reliable estimates than others.

R. Lobster Land wishes to dig just a bit deeper into its revenue trends across the 2024 season. Right now, the data is aggregated on a daily basis. Using Python, aggregate the data by week instead of day. Next, create a lineplot that shows 'Week Number' on the x-axis, and total revenue on the y-axis.

```
In [29]: # Aggregate data by week
df['Week_Number'] = df['Date'].dt.isocalendar().week
weekly_revenue = df.groupby('Week_Number')[['Total_Revenue']].sum().reset_index()

print("\nWeekly Revenue Totals:")
print(weekly_revenue)


plt.figure(figsize=(14, 6))
plt.plot(weekly_revenue['Week_Number'], weekly_revenue['Total_Revenue'],
         marker='o', linewidth=3, markersize=10, color='#D62828')
plt.xlabel('Week Number (2024)', fontsize=12, fontweight='bold')
plt.ylabel('Total Weekly Revenue ($)', fontsize=12, fontweight='bold')
plt.title('Weekly Revenue Trends - Lobster Land 2024 Season\n(Memorial Day Excluded)', fontsize=14, fontweight='bold', pad=20)
plt.grid(True, alpha=0.3)

# Add value labels for peak weeks
max_week = weekly_revenue.loc[weekly_revenue['Total_Revenue'].idxmax()]
plt.annotate(f'Peak: ${max_week["Total_Revenue"]:.0f}\n(Week {int(max_week["Week_Number"])}',
            xy=(max_week['Week_Number'], max_week['Total_Revenue']),
            xytext=(10, 20), textcoords='offset points',
            bbox=dict(boxstyle='round', pad=0.5, facecolor='yellow', alpha=0.8),
            arrowprops=dict(arrowstyle='->', connectionstyle='arc3', rad=0))

plt.tight_layout()
plt.show()
```

Weekly Revenue Totals:

	Week_Number	Total_Revenue
0	22	117681.76
1	23	108911.66
2	24	121892.56
3	25	116307.36
4	26	127882.85
5	27	122083.45
6	28	93974.40
7	29	109114.93
8	30	118805.11
9	31	105168.45
10	32	113846.59
11	33	88286.22
12	34	102002.18
13	35	96710.79
14	36	12820.86



- a. What general trend does this lineplot show, if any? If you were explaining it to an outside colleague, what would you point out about the graph?

The weekly revenue graph shows that business is strongest in the middle of the summer, with a peak around Week 26, and consistently high revenue through Week 33. This suggests Lobster Land should staff more heavily and run key promotions during mid-season, when visitor numbers are naturally highest. After Week 33, revenue drops sharply, which may reflect schools reopening at the end of peak tourism. During these later weeks, Lobster Land could reduce staffing, scale back inventory, or introduce end-of-season deals to maintain engagement while managing costs.

```
In [30]: df.to_csv("SyrettaAD654_0L_Assignment_I_Summer_2025.ipynb.csv", index=False)
```

Part III. Using Tableau to Build a Dashboard

S. Bring the **lobsterland_2_0_2_4.csv** into your Tableau Public environment.

- Dataset description available in Class Discussion >> From Your Professor

T. Using a tiled layout, build a dashboard that includes any four unique visualizations of your choice. By "unique" this just means that you should not build two of the same type of plot (e.g. not more than one histogram, not more than one barplot, not more than one treemap, etc.). Give a title to each of your four plots.

In []:

U. *Write a one-paragraph description of your dashboard. Write about the plots that you made and describe your process.* You can do this in any file format, and upload it to your Assignment 1 submission.

To explore visitor behavior and revenue trends at Lobster Land, I built a dashboard using four distinct visualizations. I started with a histogram to understand daily attendance patterns. It revealed that most days drew between 1,000 and 2,600 visitors, with over half performing above the 80% benchmark. Next, I created a line chart to track monthly revenue, which showed a dramatic dip in June, followed by a steady decline through September. Curious about external factors, I added a chart comparing revenue across weather conditions, where sunny days stood out with nearly \$500,000 in earnings. Finally, I built a table to examine the impact of special events, showing that each event type, Fireworks and Arcade Tournaments, attracted over 8,000 total visitors across the season, while regular days without events made up nearly 80% of total attendance. Throughout the process, I used a tiled layout, added reference lines, and wrote annotations to make insights easy to interpret. This dashboard tells a clear story, Lobster Land's performance is shaped by seasonality, weather, and event timing. Visualizing these patterns helped me uncover what drives the busiest and most profitable days.

V. Paste a link to your file in the same document that you used to write the description.

Note: This section is intentionally very open-ended. Each submission will be unique. The goal here is not to arrive at a single "correct" answer but to have everyone gain some hands-on experience with building a dashboard in Tableau. The dashboards will not be scored by some 'beauty contest' measure -- the goal here is to (1) make a good-faith effort to build a dashboard with four separate types of visualizations, and (2) include a thoughtful narrative paragraph. Every answer that does those things will receive credit for this section.

https://public.tableau.com/views/SyrettaAssignment1_2_0_2_5/Dashboard1?:language=en-US&:sid=&:redirect=auth&:display_count=n&:origin=viz_share_link